

High Performance Computing, Assignment 2

UID: 1690550

March 15, 2017

1 Initial readings, before any changes were made

After a clean make and running the code for the first time, we get the following results from gprof:

function	Runtime of functions (s)
poisson	23.14
compute_tentative_velocity	1.10
compute_rhs	0.12
apply_boundary_conditions	0.02
set_timestep_interval	0.02
update_velocity	0.02

Table 1: Original runtime of functions

We can see that the poisson function is by far the most expensive one, and so we shall focus on it first.

2 MPI approach

2.1 Domain Decomposition

We have the option to decompose the region in 1 or 2 directions. After some trial and error, we find that attempting to split the region both horizontally and vertically is not very helpful. It creates a lot of problems for communication between processes later, and generally we do not get much improvement this way. Additionally we can observe that the data is contiguous in columns, but not in rows.




Figure 1: Domain decomposition

Therefore we split the region into vertical chunks, one for each process, as shown in Fig. 1. It is straightforward to calculate the number of chunks necessary. We can find the number of processes

using:

```
int nprocs = 0
MPI_Comm_size(MPI_COMM_WORLD, &nprocs);
```

We put that right after initializing MPI and working out number of processes and assigning task ids, just above the main loop. For ease of calculations, we allow only a number of processes by which the size of the array is divisible. Otherwise we would get uneven chunks.

Each process is assigned a chunk consisting of $interval_size = \frac{imax}{nprocs}$ columns to balance the load.

We had to define new variables, `imin` and `iend` to denote the first and the last column of a chunk respectively. They were added as parameters to each function in 'simulation.c' (and its header).

2.2 Passing boundaries

The processes have to exchange boundaries with their neighbours. It is done by passing the column `imin` to the neighbour process on the left, and `iend` to the neighbour on the right. Of course the 1st and last process have only a single neighbour, so to exchange boundaries we use the following:

```
if (proc == 0) {
    MPI_Send(&p[iend][0], jmax+2, MPI_FLOAT, proc+1, 0, MPI_COMM_WORLD);
    MPI_Recv(&p[iend+1][0], jmax+2, MPI_FLOAT, proc+1, 1, MPI_COMM_WORLD);
}
if (proc != nprocs -1) {
    MPI_Send(&p[iend][0], jmax+2, MPI_FLOAT, proc+1, 0, MPI_COMM_WORLD);
    MPI_Recv(&p[iend+1][0], jmax+2, MPI_FLOAT, proc+1, 1, MPI_COMM_WORLD);
}...
```

Similarly for a process to the left we send column `imin`, and receive `imin-1`, and for the processes that aren't 0 or `nprocs-1` we do both. While it could be possible to have just 2 statements: if a process isn't 0 else if a process isn't `nprocs-1`, we found that splitting it into three statements was easier to debug. This is pictured in Fig. 2.




Figure 2: Boundary sharing

2.3 Reductions

There are a few places where reductions need to take place, and they are easy to miss. In each case we create a copy of the variable (naming convention: `var_global`), to reduce to.

- In `poisson()` function: `p0` is the sum of squares of the values in `p`. We need `MPI_SUM`.

- In `poisson()` function: `*res` is a more complicated function, so we create a temporary variable and a global version of the latter. Later we can calculate `*res` using `temp_global`
- In `set_timestamp_interval()` function, both `umax` and `vmax` can be reduced with `MPI_MAX`.

2.4 Gathering the data

After the main loop, we need to gather matrices `p`, `u`, and `v`, which are then printed to a file by the MASTER process. It is most efficient to do it in place. We also put an `MPI_Barrier` before gathering, to ensure that the processes have finished their calculations.

We elected to use `MPI_Gatherv`. It requires defining our data pattern. Since we want to gather it on the MASTER process (process 0), we stop it from sending anything to itself, by defining first data chunk to be of size 0 and the following ones of $(\text{number of columns in each chunk}) \times (\text{jmax}+2)$. The displacement in each case is just the size of the chunk after the previous one.

To avoid creating a copy of the array, which would be slower and memory inefficient, we use `MPI_IN_PLACE` as the parameter for the receiving buffer.

Other than the main loop, there is one more place where gathering data has to happen - at the end of `update_velocity()` function, output of which is used to apply boundary conditions.

Before each `MPI_Gatherv` statement, we put a barrier to make sure that all processes have finished their calculations by then.

2.5 Last tweaks

After all of the above work we got it to a point where it ran and the output (`karman.ppm`) was visually similar to the original, although the `.bin` files differed. It took 4.27 seconds to run the `poisson()` function on this version.

In an attempt to fix that, we have made sure that the left and right boundaries are changed only once in the `compute_tentative_velocity()` function. To this end we added a flag. The first process to execute this function in each iteration changes the flag after updating the vertical boundaries, so that other processes will not do it. However, that still was not a solution, and since it increased the runtime to about 9 s, we removed it.

We suspect that it has something to do with the rightmost boundary.

2.6 Results

Unfortunately due to the limited resources and length of the queue, we weren't able to take readings throughout the process, and the development was happening on much fewer nodes and tasks per node, so it would be difficult to compare.

We did however get a few reading on 2 nodes and 32 processes, and the final average is 8.2 s (although we did see some better times during the tests, and it seems to vary from 1.8 to 8.2, with that being the upper limit).

function	run time (s)
<code>poisson</code>	8.2
<code>update_velocity</code>	0.4
<code>compute_tentative_velocity</code>	0.02
<code>compute_rhs</code>	0.02

Table 2: Runtime of function in the final version (note that the missing functions simply did not appear in the `gprof` file, due to how quick they were)

And the visualisation of the final solution is shown in Fig. 3.



Figure 3: Visualisation of the final solution

3 Using OpenMP

First we observe that in the `simulation.c` there are multiple loops, however not all of them may be worth parallelising. We do not expect to see much improvement in performance, knowing from previous assignment that OMP's pragmas have a lot of overhead.

With the single pragma, we managed to get the average runtime of the program to be around 2 s. The solution is visually similar to the original, and does not differ from the one obtained with MPI code. However, this runtime is still possible to achieve with pure MPI approach. On the other hand the hybrid approach seems to be more consistent with running for less than 3 seconds.




Figure 4: Visualisation of the final solution with OpenMP

We have also attempted to add a pragma to the loops in the `poisson()` function, as this is still the most costly one. However we had no success, as OpenMP had issues with the `continue` statement in it. The pragma should be added inside the loop iterating over `rb`, and just above the one iterating over `i`.