CS909: Natural Language Processing, Exercise 1

UID: 1690550

March 1, 2017

1 Part B

Operation costs:

- insertion: 1

- deletion: 1

- substitution: 2

1.1 Minimum edit distance from "refa" to "fear"

	#	f	e	a	r
#	0	1	2	3	4
r	1	2	3	4	3
e	2	3	2	3	4
f	3	2	3	4	5
a	4	3	4	3	4

Table 1: Minimum edit distance of "refa" to "fear" is $\underline{4}$.

There are 2 alignments, discovered by inspection. The alignments are:

```
*   *   f   e   a   r
|   |   |   |   |   |
r   e   f   *   a   *
```

```
*   f   e   *   a   r
|   |   |   |   |   |
r   *   e   f   a   *
```

1.2 Minimum edit distance of "drive" to "brief and "drive" to "divers"

	#	b	r	i	e	f
#	0	1	2	3	4	5
d	1	2	3	4	5	6
r	2	3	2	3	4	5
i	3	4	3	2	3	4
v	4	5	4	3	4	5
e	5	6	5	4	3	4

Table 2: Minimum edit distance of "drive" to "brief" is 4

	#	d	i	v	e	r	s
#	0	1	2	3	4	5	6
d	1	0	1	2	3	4	5
r	2	1	2	3	4	3	4
i	3	2	1	2	3	4	5
v	4	3	2	1	2	3	4
e	5	4	3	2	1	2	3

Table 3: Minimum edit distance of "drive" to "divers" is 3

2   Part C

For the purpose of this problem, we define the following sets:

$$vowel = \{a, e, o, u, i\}$$

$$consonant = \Sigma - vowel$$

The set of consonants will be abbreviated to "cons.".

We are interested in an FST for the consonant doubling rule for verbs. It is important to stress that the FSTs presented below work only for regular verbs.

If the input is given on the lexical level, we can turn it into intermediate with the following FST:
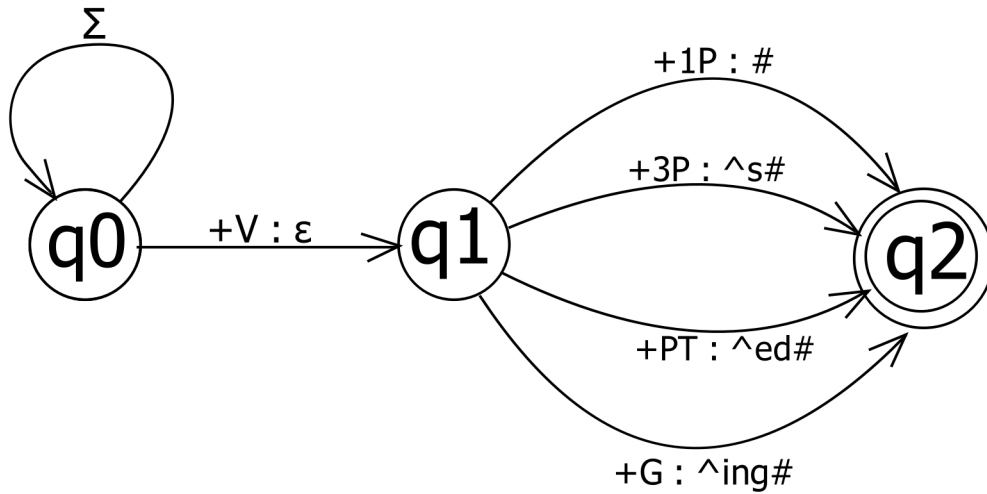

Figure 1: FST between lexical and intermediate level

Where Σ is any letter, '+1P' indicates first person ending (same as stem), '+3P' indicates third person form (i.e. -s suffix), '+PT' stands for past tense (-ed), and '+G' for gerund (-ing).

Figure 2 shows the simplest FST that encodes the consonant doubling rule just that for single syllable verbs, taking intermediate level input.
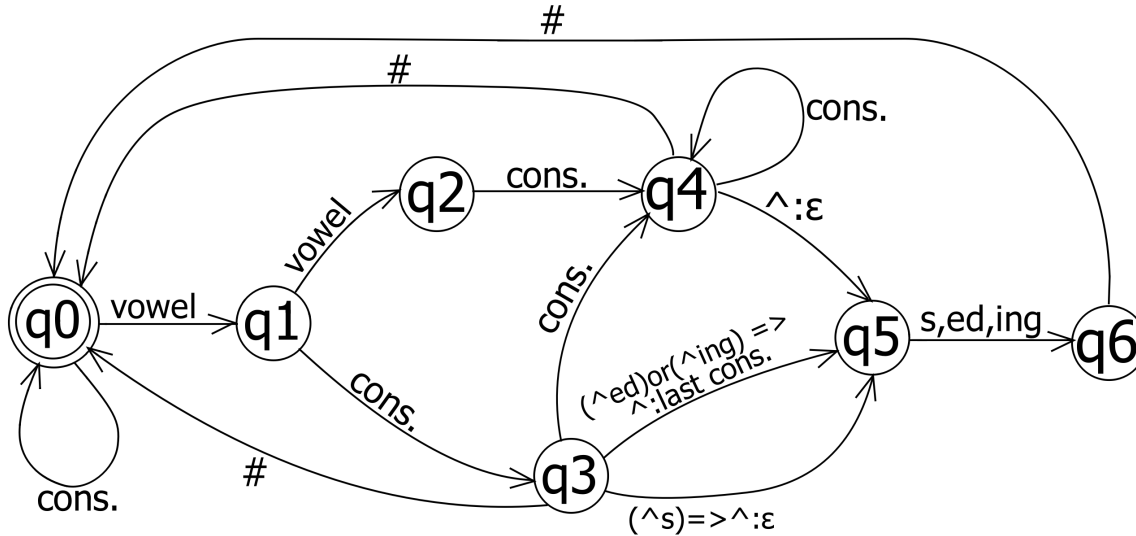

Figure 2: Simple version of the requested FST

The last consonant of the given verb will be doubled only if it is preceded by a single vowel, and if we are aiming for the regular past tense or the gerund of the verb.

Thus for example:

- "aim" will go from state q_0 to q_1, then as we are seeing a second vowel, we go to q_4, and then either we see the end of the string symbol and go to the accepting state q_0, or add a required suffix.

- "bring" will loop over the "cons." edge around q_0 twice, then go to q_1 and on to q_3. Then, since we're seeing another consonant, we move to q_4 and avoid doubling the last letter.

- "stop" or "fit" will loop around q_0 until they get to q_1 and q_3. If the next character is the end of string symbol, we move back to the accepting state q_0. Otherwise, we read ahead to see what ending we need. If it's "s", we replace "∧" with empty symbol. Otherwise, we replace it with the last consonant.

However, this FST does not treat "e" at the end of the words correctly. Therefore we present a more robust version, which works for regular verbs, including those with multiple syllables and those ending in e (silent or otherwise).

It is perhaps surprising that only states directly concerned with the consonants doubling are q_0, q_1, q_2, q_6, q_7, and q_{10}. The state q_5 is also important, as together with q_1 and q_2 it helps with detecting the 'consonant - vowel - consonant - ∧' pattern. The rest deals with rules regarding e at the end of words.

To explain what it does at each step, we shall inspect all outgoing edges from each node.

- q_0: we are coming back to it until we see the first consonant in the verb. Then we move to q_1
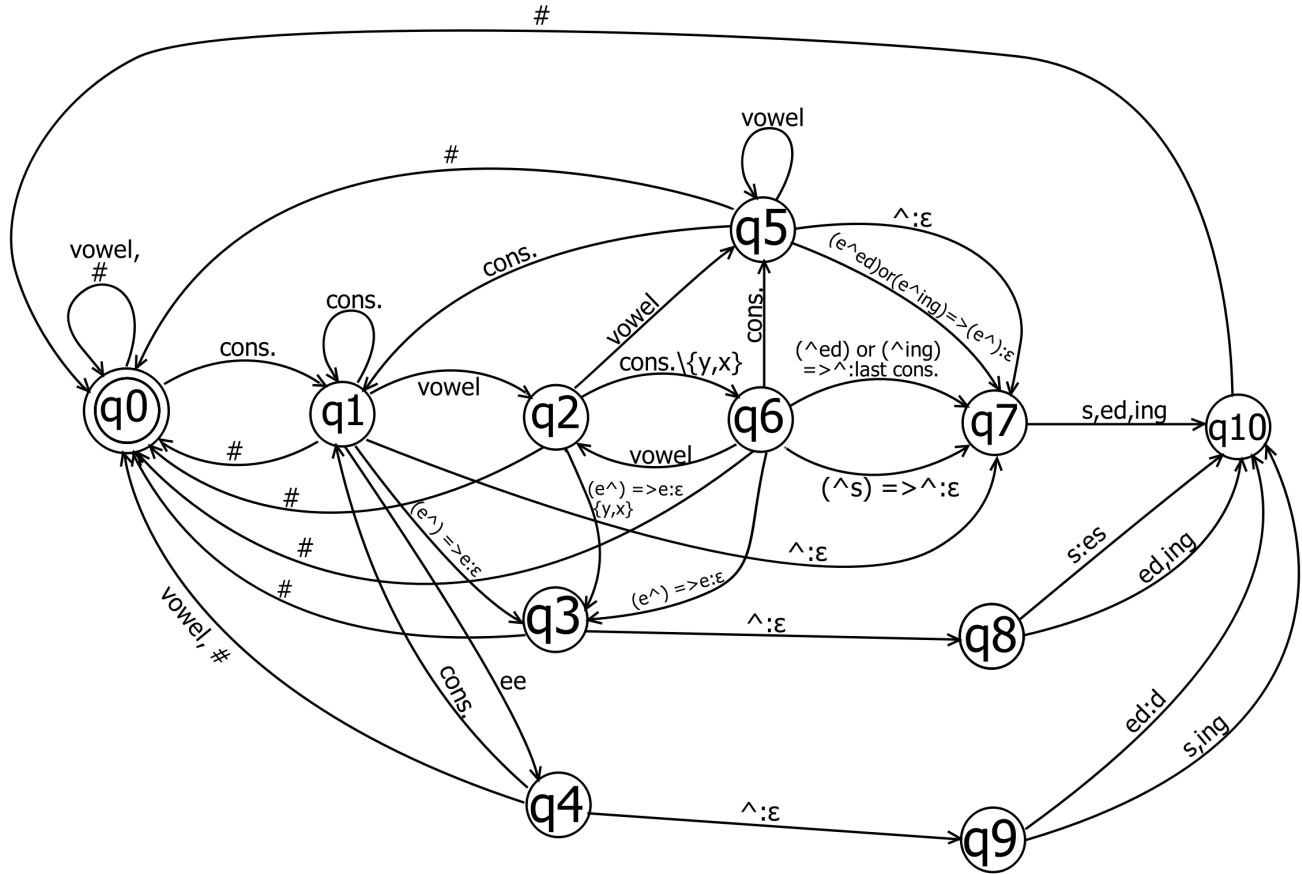
3

Figure 3: More robust FST

- q_1: this is the state from which be begin counting the consonants. I.e. we get to it after following a consonant. If we then see a vowel and a consonant, we apply the doubling rule. Other options are: double e (in words such as "agree") or a silent e where we focus on dealing with the rules regarding this, or "∧".

- q_2: the last consonant in a word here will get doubled, and so we proceed to q_6. However, the exceptions to this rule are the letters x and y. In these cases as well as a word-ending e (e.g. in "woe") we move to q_3. The last option is any vowel, leading to q_5.

- q_3: this is the state which catches "e ∧" at the end of the stem. Note that we cannot get to q_3 without seeing ∧, so we do not need an outgoing edge marked with end of string # character.

- q_4: we treat double e separately, because in the case of past tense we lose the e in the suffix, but otherwise we don't alter the stem. There is a very subtle difference between states q_8 and q_9 which we will analyse later. Words such as "agree" or "need" will go through this state.

- q_5: this state is part of the cycle detecting 'consonant - vowel - consonant' pattern. It will also detect the technically regular but difficult words such as "queue".[1]

- q_6: We can only get to this state after seeing the desired "consonant - vowel - consonant" pattern, and indeed all the incoming edges are from vowels. The two most important outgoing

---

[1]It is very likely that some simplifications can be made around this state

edges describe the consonant doubling rule and lead to q_7. The others are useful for when we have not reached the end of the word after spotting the pattern.

- q_7, q_8, and q_9 all deal with adding the desired suffix. In the first case, we take care of verbs ending with consonants (doubled or not). The latter two deal with words ending with e. Most of the words ending with a silent e will end up in q_8. To simplify dealing with it, we have removed the silent e, and so we can simply append the '-ed' and '-ing' suffixes. For the '-s' suffix, we need to add the e back again. q_9 is concerned with words ending in double e, such as 'agree' or 'free'. In these cases we did not remove any of the e's previously, and we can simply append '-ing' and '-s', but '-ed' gets shortened to just '-d'.

- Finally q_{10} is reached after appending the desired suffix, and the only character we expect to see here is the end of string character, so we send the input to the accepting state q_0.

Here are some example of words accepted by this FST:

- "aim#" (aim): $q_0 \to q_0 \to q_0 \to q_1 \to q_0$

- "aim$\wedge$ing#" (aiming): $q_0 \to q_0 \to q_0 \to q_1 \to q_7 \to q_{10} \to q_0$

- "stop$\wedge$s#" (stops): $q_0 \to q_1 \to q_1 \to q_2 \to q_6 \to$ (following the bottom edge, no doubling) $q_7 \to q_{10} \to q_0$

- "stop$\wedge$ed#" (stopped): $q_0 \to q_1 \to q_1 \to q_2 \to q_6 \to$ (following the top edge, doubling) $q_7 \to q_{10} \to q_0$

- "free$\wedge$ed#" (freed): $q_0 \to q_1 \to q_1 \to q_4 \to q_9 \to$ (following the top edge, to avoid triple e) $q_{10} \to q_0$

- "like$\wedge$ing#" (liking): $q_0 \to q_1 \to q_2 \to q_6 \to q_3 \to q_8 \to$ (following the bottom edge) $\to q_{10} \to q_0$

- "queue$\wedge$ed#" (queued): $q_0 \to q_1 \to q_2 \to q_5 \to q_5 \to q_7 \to q_{10} \to q_0$

3 Part D

We are given the corpus:

```
<s> I am Sam </s>
<s> Sam I am </s>
<s> I am Sam </s>
<s> I do not like green eggs and Sam </s>
```

We see the word 'Sam' 4 times, twice following 'am', and the word 'am' 3 times. With $\lambda_1 = \dfrac{1}{2}$ and $\lambda_2 = \dfrac{1}{2}$, the linear interpolation smoothing for unigrams and bigrams is given by:

$$\hat{\mathbb{P}}(Sam|am) = \lambda_1 \mathbb{P}(Sam|am) + \lambda_2 \mathbb{P}(Sam)$$

$$= \frac{1}{2} \cdot \frac{2}{3} + \frac{1}{2} \cdot 4 = \frac{1}{3} + 2 = 2\frac{1}{3} \approx 2.33$$